Transfer Learning for Machine Comprehension over Bi-Directional Attention Flow Networks

Carlos Castro

`carlosscastro@berkeley.edu`

Abstract

Machine text comprehension involves answering a query about a given context paragraph, and requires modeling complex interactions between the context and the question. In the recent years, with the release of the Stanford Question Answering dataset (SQuAD) (Rajpurkar et al., 2016) and the Microsoft MAchine Reading COmprehension Dataset (MS-MARCO) dataset (Nguyen et al., 2016), there were significant improvements in the state of the art, to the point where some neural network architectures are relatively close to achieving human level accuracy on the SQuAD dataset. Despite the great advances in the field of machine reading comprehension, training state of the art models requires humongous amounts of labeled data consisting of thousands of passages plus multiple question-answer pairs on each passage. This acts as a limiting factor in the applicability of machine comprehension techniques in technology outside of research benchmarks. In this paper, we aim to provide a solution to this shortcoming, by studying transfer learning over a neural network trained using the SQuAD dataset to other corpora. Successful results would open the door to multiple direct applications and further research opportunities.

1 Introduction

Since the Stanford Question Answering Dataset (SQuAD) (Rajpurkar et al., 2016) was released, rapid progress was made in the field of machine question answering. The original paper already proposed a strong logistic regression model, and later other better performing approaches were published, based off match-LSTMs (Wang et al.) and bi-directional attention flow networks (Seo et al., 2017). Despite this recent progress in the field of machine text comprehension, SQuAD is a luxury dataset with a high amount of labeled data. This makes the applicability of these approaches to other corpora somewhat limited, since such humongous training data is not available.

In the field of computer vision, neural networks are rarely trained from scratch. Generally networks are trained with ImageNet (Deng et al.), a large-scale hierarchical image database, to obtain features. Transfer learning allows us to transfer this knowledge to other tasks. Analogously to ImageNet for computer vision, we could use weights from neural networks trained for SQuAD dataset as initial weights for different tasks in other corpora.

In this paper, we study training bi-directional attention flow networks (BiDAF) on the SQuAD dataset and then transferring that knowledge to perform on the Microsoft MAchine Reading COmprehension Dataset (MS-MARCO) dataset (Nguyen et al., 2016), which is based off real web queries and human written answers. Given that the SQuAD dataset is extracted from Wikipedia content (Rajpurkar et al., 2016), while the MS-MARCO dataset is obtained from web user queries and web content, the domains and quality of the text vary wildly. With MS-MARCO being quite representative of web content, it contains a wide variety of text and content quality, while the Wikipedia content from the SQuAD dataset is quite curated and edited by the community. Considering this, transfer learning from a model trained on the SQuAD dataset to the MS-MARCO dataset is a non-trivial task, and with relevant implications as we discuss in sub-section 1.1.

1.1 Motivation

The implications of successful transfer learning from models trained on the SQuAD dataset to other corpora with little or no labeled data are extremely relevant for most conversational interfaces, including chatbots, conversational agents and conversation-augmented applications. Below, we describe two brief scenarios that demonstrate potential applicability of this study.

First, consider question-answering chatbots. It might be useful to have chatbots that receive certain passages and can answer questions about those passages. A concrete example would be a car-embedded chatbot that ingests the car manual, and then the passengers can ask questions to the car, such as *what is the ideal tire pressure?*. Should transfer learning be proven feasible over BiDAF networks, this chatbot could be created with little or none labeled examples.

Another example could be augmenting web content with conversational agents. A possible implementation could be a web browser extension to which users can ask questions about the current web page. A user loads a web page and asks a question, then the text and the question are entered into a model trained over the SQuAD dataset, and the user gets an answer to their query without even reading the article.

Note that without the use of transfer learning, it would not be possible to enable the scenarios described above, because this is such a complex problem space that a rather large number of labeled documents is required to train a model. This is supported by the fact that machine comprehension state of the art skyrocketed after the release of the SQuAD dataset (Natural Language Computing Group) (Seo et al., 2017).

1.2 Paper organization

In section 2 we discuss bibliography and related work around machine comprehension datasets, transfer learning and machine comprehension neural network architectures. Section 3 describes the approach and experiments to study this problem space, and later we review the experiment results in section 4. We outline the next steps towards this study in section 6.

Passage: Tesla later approached Morgan to ask for more funds to build a more powerful transmitter. When asked where all the money had gone, Tesla responded by saying that he was affected by the Panic of 1901, which he (Morgan) had caused. Morgan was shocked by the reminder of his part in the stock market crash and by Teslas breach of contract by asking for more funds. Tesla wrote another plea to Morgan, but it was also fruitless. Morgan still owed Tesla money on the original agreement, and Tesla had been facing foreclosure even before construction of the tower began.
Question: On what did Tesla blame for the loss of the initial money?
Answer: Panic of 1901

Table 1: An example from the SQuAD dataset.

2 Related work

2.1 SQuAD Dataset

The Stanford Question Answering Dataset (SQuAD) (Natural Language Computing Group), is a reading comprehension dataset consisting of 100,000+ questions posed by crowd-workers on a set of Wikipedia articles, where the answer to each question is a segment of text from the corresponding reading passage (Rajpurkar et al., 2016).

In the SQuAD dataset, the answer to every question is a segment of text, or *span*, from the corresponding reading passage. SQuAD contains 107,785 question-answer pairs on 536 Wikipedia articles.

2.2 MS-MARCO Dataset

The Microsoft MAchine Reading COmprehension (MS-MARCO) (Nguyen et al., 2016) dataset is aimed to overcome a number of well-known weaknesses of previous publicly available datasets for the same task of reading comprehension and question answering. In MS-MARCO, all questions are sampled from real anonymized user queries. The context passages, from which answers in the dataset are derived, are extracted from real web documents using the most advanced version of the Bing search engine. The answers to the queries are human generated, and may contain words that could not be in the original paragraph. This last consideration is a crucial difference with the SQuAD dataset. Additionally, while the SQuAD

Passage: The goal you choose will determine your path. A clinical psychologist, for example, will have heavy training in both the theory and practice of psychology. Youll typically need a doctorate degree in a psychology-related field in order to build a career. A psychiatrist has even more rigorous demands; becoming one requires medical school training. In order to do this, you have to do three things: 1) Get work experience under the supervision of a licensed professional (usually for two years); 2) Pass the board exams; and 3) Depending on your state, present a valid case study to the board. After all that work, youll finally be able to call yourself a psychologist!
Question: do you have to do a phd to be a clinical psychologist
Answer: Yes

Table 2: An example from the MS-MARCO dataset.

dataset contains Wikipedia content which is generally quite curated, the MS-MARCO dataset contains a widea breadth of writing styles, such as news, blogs, social media, etc. MS-MARCO contains 100,000 queries with their corresponding answers.

For the purpose of this study, we use only questions from the MS-MARCO dataset where the answer is a sub-span of the passage. This way we focus on studying the implications of answering questions from a different domain, without introducing yet the additional complexity of dealing with words outside of the original question.

2.3 Reading comprehension

For reading comprehension style question answering, a passage **P** and question **Q** are given, our task is to predict an answer **A** to question **Q** based on information found in **P**. The SQuAD dataset further constrains answer **A** to be a continuous sub-span of passage **P**. Answer **A** often includes non-entities and can be much longer phrases. This setup challenges us to understand and reason about both the question and passage in order to infer the answer. Table 1 shows a simple example from the SQuAD dataset. As for MS-MARCO dataset, several related passages **P** from Bing Index are provided for a question **Q**. Besides, the answer A in MS-MARCO is generated by human which can

not be a continuous sub-span of the passage.

2.4 Bi-Directional Attention Flow

To study transfer learning, we choose one method out of the multiple state of the art techniques, named Bi-Directional Attention Flow (BiDAF) network (Seo et al., 2017). BiDAF is a hierarchical, multi-stage architecture for modeling the representations of the context paragraph at different levels of granularity. BIDAF includes character-level, word-level, and contextual embeddings, and uses bi-directional attention flow to obtain a query-aware context representation. It would be interesting studying how other state of the art techniques, such as R-NET, respond to transfer learning.

The implementation of Bi-directional attention flow networks we utilize (Seo et al., 2017) is specially designed for datasets where the anwer A is a sub string of the passage P. This is clearly reflected in the architecture of the network, where the inputs are the context passage and the query, and output are basically two integers, indicating the start and end indices of the answer text within the context passage. To achieve this, the training loss to be minimized can be expressed as the sum of the negative log probabilities of the true start and end indices by the predicted distributions, averaged over all examples:

$$L(\theta) = -\frac{1}{N} \sum_{i}^{N} log(p_{y_i^1}^1) + log(p_{y_i^2}^2)$$

where θ is the set of all trainable weights in the model including CNN filters and LSTM cells, N is the number of train examples and y_i^1 and y_i^2 are the true start and end indices of the i-th example respectively and p_k indicates the k-th value of the vector p, which represents the probability distributions of the start and end indices over the entire paragraphs, and is obtained from the softmax calculation in the output layer (Seo et al., 2017).

2.5 Transfer learning

Here we provide useful definitions and conventions related to transfer learning. Let a domain **D** consist of a feature space **X** and a marfinal probability distribution $P(X)$. Given a source domain **D**$_s$ and a learning task **T**$_s$, a target domain **D**$_t$ and a target task **T**$_t$, *trasnfer learning* aims to improve
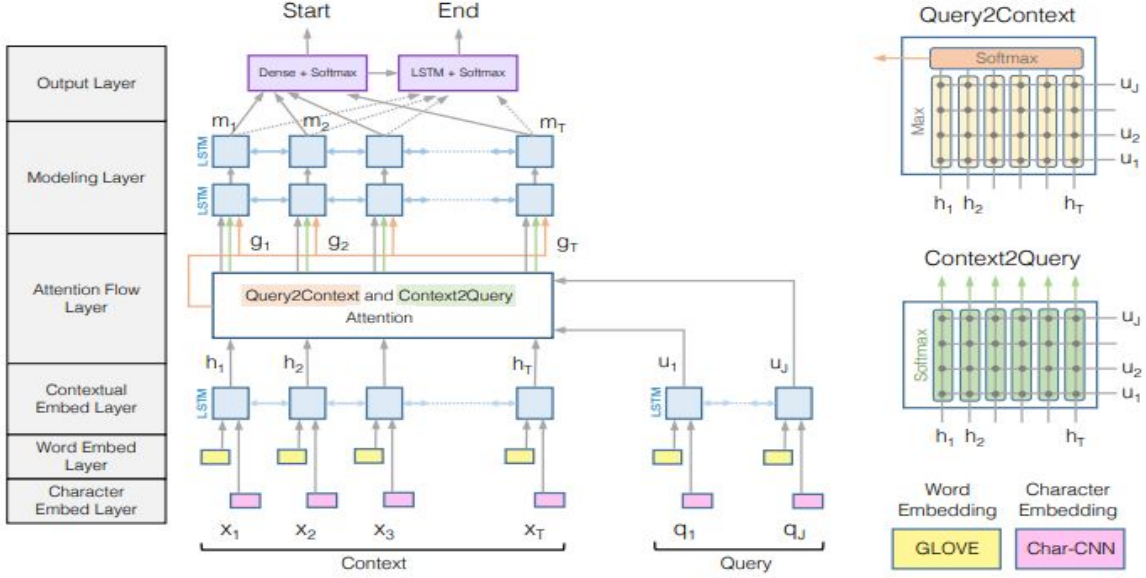
Figure 1: Bi-directional attention flow network architecture

the learning of the predictive function for $\mathbf{D}_t$ using the knowledge in $\mathbf{D}_s$ and $\mathbf{T}_s$. In the *inductive transfer learning* setting, the target task is different from the source task, no matter when the source and target domains are the same of not. Conversely, in *transductive transfer learning*, the source and target tasks are the same, while the domains are different. In this situation, little or no labeled data are available in the target domain, while lots of data are present for the source domain (Pan and Yang) (Conneau et al., 2017) (Peters et al.) (Conneau et al., 2017).

In computer vision, neural network models are rarely trained from scratch. In general, initial weights are the result of training with ImageNet (Deng et al.), a large-scale hierarchical image database, to obtain features. There is already relevant work studying transfer learning over text tasks, mostly around vector representations of words.

3 Methods

In this section, we describe our methods to study transfer learning over BiDAF networks.

First, we train a BiDAF network on the SQuAD dataset, being this our source task and domain, and we'll refer to the resulting model as the *source model*. Since we want to study how we can transfer the knowledge to other tasks and domains with none or limited labeled data, we perform a number of transfer experiments with varying amounts

of labeled question-answer pairs from the MS-MARCO dataset, which will be our target task and domain for transfer learning.

Let n be the number of labeled samples in the target domain we want to experiment with, then an experiment is as follows: Select n random question-answer pairs from the MS-MARCO dataset, re-train the source model on the randomly selected pairs, and evaluate the resulting model on the MS-MARCO evaluation set. We run this experiment multiple times, for different values of n, including $n = 0$, which is the case where the target domain has 0 labeled examples.

Finally, we also train another BiDAF network on the MS-MARCO dataset without transfer, as baseline that represents the performance we would reach without transfer learning, given a very small dataset. The ideal result of this experiment would be to achieve significantly higher accuracy in our transfer experiments than on this non-transfer model, which would mean that even in domains with little data, transfer learning from SQuAD can make up for that lack of data. sIt is worth noting that BiDAF performance over MS-MARCO was never studied, so that is an unexpected contribution of our work, in addition to its main focus around transfer learning. We won't report training over the entire marco dataset, but just a subset. We will report the results on the full dataset in a future separate study.

4 Results

4.1 Bi-DAF over SQuAD dataset

The first step of our research is to train a BiDAF over the SQuAD dataset. We train for 16.4 hours using a single NVidia K80 GPU, and obtain results very close to those reported in the original study (Seo et al., 2017), with exact match percentage of 65.64 and F1 score of 75.68. The baseline logistic regression performance for this dataset is exact match percentage of 40.0 and F1 score of 51. Human performance for this task is EM at 82.3 and F1 score of 91.2 (Rajpurkar et al., 2016) (Natural Language Computing Group).

Note that as most literature around this style of machine comprehension, we report EM which represents the percentage of exact matches, and the F1 score, which measures the weighted average of the precision and recall rate at a character level.

4.2 Bi-DAF over MS-MARCO dataset

We start by obtaining 500 labeled examples from the MS-MARCO dataset for training and 250 for testing, all in which the answer is a substring of the context passage. We restrict the number of samples, since the purpose of this study is to understand the applicability of architectures such as Bi-DAF to other domains where we have little to no labeled data.

Once we have the labeled samples from MS-MARCO, we run 2 families of experiments: first, we study direct training a new Bi-DAF over that data and second, we study fine tuning Bi-DAF networks pre-trained over SQuAD dataset with two different learning rates.

For the first experiment, we train a Bi-DAF from scratch over the MS-MARCO 500 labeled samples with no transfer. Given the little amount of data, we expect poor results, but this allows us to compare with the results we'll obtain later from transfer learning, and emphasize the value of knowledge transfer. Many studies on transfer learning start by using direct training transfer path as a baseline or means of comparison (Conneau et al., 2017).

For the second set of experiments, we start with two Bi-DAF networks trained with the SQuAD dataset, using initial learning rates of 0.8 and 0.5 respectively. The motivation for this, is that even though the lower learning rate is sub optimal when verifying for SQuAD, it might help to avoid overfitting to the SQuAD domain, and allow for better adaption of the network to the MS-MARCO content. Before fine tuning the network with our MS-MARCO samples, we evaluate the SQuAD trained network over the MS-MARCO evaluation set, to also understand how much the fine tuning does later on. We incrementally fine-tune the pre-trained model with different amount of samples until we reach the total of 500 training samples, and evaluate against our MS-MARCO training samples at each step, to understand the impact of the number of samples used for fine tuning in the transfer performance.

The results over the MS-MARCO evaluation is depicted in the table below.

5 Analysis

5.1 Transfer Learning value

There are many interesting things exposed by the results. First, the low performance when directly training Bi-DAF network over a dataset with a low number of samples is evident as we obtained an F1 score of 11.02. This highlights the need for transfer learning towards other domains, since we can't generally obtain domain-specific and rich datasets such as SQuAD. Second, the substantial progress that can be achieved over a bit of fine tuning using a handful of domain-specific samples. With only 500 samples for fine tuning, we increased the F1 score more than 10 percent. This supports the theory in the near future, the SQuAD could be a staple dataset for neural model priming prior to transfer learning, such as ImageNet is in the field computer vision.

It is also worth noting that even though pre-training the model with lower learning rate resulted in worst performance over SQuAD, when applied to transfer learning showed better results, likely because this allowed the network to generalize and perform better when being re-trained and subsequently shown passages and queries from a different domain.

5.2 Error Analysis

Of course, an F1 score of 59.14 is still somewhat low and needs to be improved prior to start applying it in concrete technology. We now evaluate error areas to understand where we are failing, and identify areas for future research. We analyzed the errors for the most accurate model: Bi-DAF SQuAD pre-trained + MS-MARCO fine tuning (500 samples) with initial learning rate of

Model	EM	F1
Logistic regression baseline (Rajpurkar et al., 2016)	40.4	51.0
Bi-DAF (Original paper)	68.0	77.3
Bi-DAF (Our result - initial learning rate: 0.8)	66.905	75.68
Bi-DAF (Our result - initial learning rate: 0.5)	65.71	74.07
R-Net (Single model)	76.461	84.265
Human performance (Rajpurkar et al., 2016)	82.304	91.221

Table 3: Results over the SQuAD dataset, including our results, baseline, human performance and state of the art as of the time of writing this work. All results are for single model, ensembles tend to do quite better but add complexity to the study.

Model	F1 (lr = 0.8)	F1 (lr = 0.5)
Bi-DAF over MS-MARCO baseline (Direct training - no transfer)	11.02	n/a
Bi-DAF SQuAD pre-trained (direct evaluation - no transfer)	50.97	48.33
Bi-DAF SQuAD pre-trained + MS-MARCO fine tuning (133 samples)	54.88	55.02
Bi-DAF SQuAD pre-trained + MS-MARCO fine tuning (250 samples)	55.71	56.83
Bi-DAF SQuAD pre-trained + MS-MARCO fine tuning (500 samples)	56.01	**59.14**

Table 4: Results for transfer learning to the MS-MARCO domain, including the no-transfer baseline and the no fine tuning baseline.

0.5. Still, F1 scores like these were the state of the art for SQuAD a when it was released, so it is still a great achievement given that we are talking about accuracy on a 500 sample dataset.

Now, we analyze the impact of the kind of answer for evaluating question performance. In the MS-MARCO subset que utilized, we had most of the questions being of numeric answer (28.4%, description answer (52.6%) and location answer (5.7%). We suspected that there might be specific types of questions where we underperformed, but that was not the case. The within-category F1 scores were never more than 3% apart from each other.

Analyzing the failing questions, we found 3 specific predominant traits: The formality and correctness of the language, the type of question (Why, when, what, etc) and whether the question has or not an answer.

First, most of the wrongly answered questions had informal writing that would never be in Wikipedia, such as text from social networks or specific informal news.

Second, we observe similar accuracy and F1 scores for questions that started with *Where* and *What* for example ($F1 = 62.02$ and $F1 = 61.875$ respectively), but a notorious decrease in performance for *Why* questions ($F1 = 51.105$). The Bi-DAF original authors also report low perfor-

mance in *Why* and *Where* questions, so we think this might be part of the nature of the architecture, given that these types of questions are as common as the other types in the original dataset, but it needs to be subjected to further analysis.

Finally, we observe very low F1 scores for questions that do not have an answer ($F1 = 0.48.12$), which means that our tended to provide an incorrect answer rather than predict that there was no answer. However, in our 500 samples dataset, there were only two such questions, so this low perofrmance might be related to imbalances in our training set and underlying dataset.

6 Future work

Overall, the goals following this study center around improving the performance of transfer learning to reach performances that allow the direct applicability of these concepts to technology. Below we mention some potential research directions.

First, as a direct continuation of this study, we will run multiple experiments to compensate for the issues found in the error analysis, such as questions without answers, *why* questions, etc.

Additionally, other model-specific techniques should be studied, such as a broader range of learning rates for the pre-training over the SQuAD dataset. Another experiment to consider would be

to not freeze the embedding layer after completing the SQuAD training, and rather keep training embeddings during fine tuning.

Another crucial direction would be to now apply this over the top performing ensembles. For simplicity, we did this study over Bi-DAF which achieves $F1 = 77.3$ over SQuAD. However, there are Bi-DAF based ensembles from the same authors in the leader board that achieve much higher results, such as the recent BiDAF + Self Attention + ELMo (ensemble) which obtained $F1 = 87.432$, and is currently in the 5th position in the SQuAD leader board. Should that performance improvement translate into our transfer experiments, we would be already obtaining a substantial increase in the overall applicability of this approach.

References

A. Conneau et al. 2017. Supervised learning of universal sentence representations from natural language inference data.

Jia Deng et al. Imagenet: A large-scale hierarchical image database.

Microsoft Research Asia Natural Language Computing Group. R-net: Machine reading comprehension with self-matching networks.

Tri Nguyen et al. 2016. Ms marco: A human generated machine reading comprehension dataset.

Sinno Jialin Pan and Qiang Yang. A survey on transfer learning.

Matthew Peters et al. Deep contextualized word representations.

Pranav Rajpurkar et al. 2016. Squad: 100,000+ questions for machine comprehension of text.

Minjoon Seo et al. 2017. Bi-directional attention flow for machine comprehension.

Wang et al. Machine comprehension using match-lstm and answer pointer.